

AWS AI Practitioner Study Guide

Domain 1.1: Fundamentals of AI and ML - Basic AI Concepts and Terminologies

1. Core AI Terminology Definitions

Artificial Intelligence (AI)

Definition: The simulation of human intelligence in machines that are programmed to think, learn, and make decisions like humans.

- **Key characteristics:** Problem-solving, reasoning, perception, learning
- **Examples:** Virtual assistants (Alexa, Siri), recommendation systems, autonomous vehicles
- **AWS Services:** Amazon Comprehend, Amazon Rekognition, Amazon Lex

Machine Learning (ML)

Definition: A subset of AI that enables computers to learn and improve from experience without being explicitly programmed for every task.

- **Core concept:** Algorithms that can identify patterns in data and make predictions
- **Process:** Train models on data → Model learns patterns → Apply model to new data
- **AWS Services:** Amazon SageMaker, Amazon Machine Learning

Deep Learning

Definition: A subset of machine learning that uses artificial neural networks with multiple layers (hence "deep") to model and understand complex patterns.

- **Architecture:** Multiple interconnected layers of neurons
- **Strengths:** Excellent for image recognition, natural language processing, speech recognition
- **Requirements:** Large amounts of data and computational power
- **AWS Services:** Amazon SageMaker (supports deep learning frameworks)

Neural Networks

Definition: Computing systems inspired by biological neural networks, consisting of interconnected nodes (neurons) that process information.

- **Components:** Input layer, hidden layers, output layer
- **Function:** Each neuron receives inputs, applies weights, and passes output through activation function
- **Types:** Feedforward, convolutional (CNN), recurrent (RNN), transformer networks

Computer Vision

Definition: AI field that enables computers to interpret and understand visual information from the world.

- **Applications:** Image classification, object detection, facial recognition, medical imaging
- **Techniques:** Convolutional neural networks, image preprocessing, feature extraction
- **AWS Services:** Amazon Rekognition, Amazon Textract

Natural Language Processing (NLP)

Definition: AI branch that helps computers understand, interpret, and generate human language.

- **Tasks:** Sentiment analysis, language translation, text summarization, chatbots
- **Challenges:** Context understanding, ambiguity, cultural nuances
- **AWS Services:** Amazon Comprehend, Amazon Translate, Amazon Polly, Amazon Transcribe

Model

Definition: A mathematical representation of a real-world process, trained on data to make predictions or decisions.

- **Components:** Parameters, weights, biases learned during training
- **Types:** Classification models, regression models, clustering models
- **Lifecycle:** Training → Validation → Testing → Deployment → Monitoring

Algorithm

Definition: A set of rules or instructions that tells the computer how to solve a problem or perform a task.

- **ML Context:** The mathematical procedure used to find patterns in data
- **Examples:** Linear regression, decision trees, random forest, support vector machines
- **Selection criteria:** Problem type, data size, accuracy requirements, interpretability needs

Training

Definition: The process of teaching a machine learning model to make predictions by showing it examples from a dataset.

- **Process:** Algorithm learns patterns by adjusting parameters based on training data
- **Components:** Training data, loss function, optimization algorithm
- **Goal:** Minimize the difference between predicted and actual outcomes

Inferencing

Definition: The process of using a trained model to make predictions on new, unseen data.

- **Also called:** Prediction, scoring, evaluation
- **Input:** New data points
- **Output:** Predictions, classifications, or recommendations
- **Performance:** Measured by accuracy, speed, and resource consumption

Bias

Definition: Systematic errors or prejudices in AI systems that can lead to unfair or inaccurate outcomes.

- **Types:**

- **Data bias:** Unrepresentative training data
- **Algorithmic bias:** Flawed model assumptions
- **Confirmation bias:** Reinforcing existing beliefs
- **Impact:** Discriminatory outcomes, reduced accuracy for certain groups
- **Mitigation:** Diverse datasets, bias testing, regular auditing

Fairness

Definition: The principle that AI systems should treat all individuals and groups equitably without discrimination.

- **Measurements:** Equal opportunity, demographic parity, individual fairness
- **Challenges:** Balancing competing fairness definitions, cultural differences
- **AWS Tools:** Amazon SageMaker Clarify for bias detection

Fit (Model Fit)

Definition: How well a model performs on data, with three main categories:

- **Underfitting:** Model is too simple, poor performance on both training and test data
 - **Symptoms:** High bias, low complexity
 - **Solutions:** Add features, increase model complexity, reduce regularization
- **Good Fit:** Model performs well on both training and test data
 - **Characteristics:** Balanced bias-variance tradeoff
 - **Goal:** Optimal generalization to new data
- **Overfitting:** Model memorizes training data but performs poorly on new data
 - **Symptoms:** High variance, too complex for the data
 - **Solutions:** More data, regularization, cross-validation, simpler models

Large Language Model (LLM)

Definition: Advanced AI models trained on vast amounts of text data to understand and generate human-like text.

- **Characteristics:** Billions of parameters, transformer architecture, pre-trained on diverse text
- **Capabilities:** Text generation, translation, summarization, question answering, code generation
- **Examples:** GPT models, BERT, Amazon Titan, Anthropic Claude
- **AWS Services:** Amazon Bedrock, Amazon CodeWhisperer

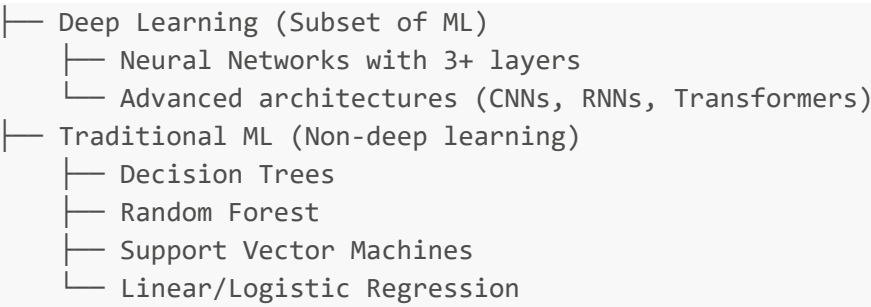
2. Relationships Between AI, ML, and Deep Learning

Hierarchical Relationship

```

AI (Broadest)
├─ Machine Learning (Subset of AI)

```



Key Similarities

- **All aim to solve problems** that traditionally required human intelligence
- **All learn from data** to improve performance over time
- **All make predictions** or decisions based on input data
- **All require training** on relevant datasets
- **All can be evaluated** using performance metrics

Key Differences

Aspect	AI	Machine Learning	Deep Learning
Scope	Broadest field	Subset of AI	Subset of ML
Data Requirements	Varies	Moderate amounts	Large amounts
Human Intervention	Can be rule-based	Learns patterns automatically	Minimal feature engineering
Computational Power	Varies	Moderate	High (GPUs/TPUs)
Interpretability	Varies	Often interpretable	Often "black box"
Examples	Rule-based systems, expert systems	Linear regression, decision trees	Neural networks, CNNs, RNNs

3. Types of Inferencing

Batch Inferencing

Definition: Processing multiple data points together in groups (batches) at scheduled intervals.

Characteristics:

- **Timing:** Asynchronous, scheduled processing
- **Volume:** Large datasets processed together
- **Latency:** Higher latency acceptable (minutes to hours)
- **Cost:** More cost-effective for large volumes
- **Use cases:** Monthly reports, data warehouse updates, bulk image processing

AWS Implementation:

- Amazon SageMaker Batch Transform
- AWS Batch for large-scale processing
- Amazon EMR for distributed batch processing

Example Scenario: A retail company processes all customer transaction data overnight to generate personalized product recommendations for the next day.

Real-time Inferencing

Definition: Processing individual data points immediately as they arrive, providing instant responses.

Characteristics:

- **Timing:** Synchronous, immediate processing
- **Volume:** Individual or small batches
- **Latency:** Low latency required (milliseconds to seconds)
- **Cost:** Higher per-request cost
- **Use cases:** Fraud detection, recommendation engines, chatbots, autonomous vehicles

AWS Implementation:

- Amazon SageMaker Real-time Endpoints
- Amazon API Gateway + AWS Lambda
- Amazon Kinesis for streaming data

Example Scenario: A credit card company analyzes each transaction in real-time to detect potential fraud and either approve or decline the transaction instantly.

Comparison Table

Factor	Batch Inferencing	Real-time Inferencing
Response Time	Minutes to hours	Milliseconds to seconds
Data Volume	Large datasets	Individual records
Resource Usage	Efficient for bulk	Higher per-request overhead
Complexity	Simpler architecture	Requires scalable infrastructure
Cost	Lower per prediction	Higher per prediction
Scheduling	Planned intervals	On-demand

4. Types of Data in AI Models

Labeled vs. Unlabeled Data

Labeled Data

Definition: Data that includes both input features and known correct outputs (ground truth).

- **Structure:** Input-output pairs
- **Examples:**
 - Email marked as "spam" or "not spam"
 - Images tagged with object names
 - Customer reviews with sentiment ratings
- **Use:** Supervised learning algorithms
- **Challenges:** Expensive and time-consuming to create
- **Quality:** Accuracy of labels directly impacts model performance

Unlabeled Data

Definition: Data that contains only input features without corresponding correct outputs.

- **Structure:** Input features only
- **Examples:**
 - Raw text documents without categories
 - Images without identification tags
 - Customer transaction records without fraud labels
- **Use:** Unsupervised learning, semi-supervised learning
- **Advantages:** Abundant and less expensive to collect
- **Challenges:** Harder to validate model performance

Data Structure Types

Tabular Data

Definition: Data organized in rows and columns, similar to spreadsheets or databases.

- **Structure:** Structured format with clear relationships
- **Examples:** Customer demographics, sales records, financial data
- **Features:** Each column represents a feature/attribute
- **Processing:** Traditional ML algorithms work well (decision trees, linear regression)
- **AWS Storage:** Amazon RDS, Amazon Redshift, Amazon S3 (CSV, Parquet)

Time-series Data

Definition: Data points collected or recorded at specific time intervals, showing how values change over time.

- **Characteristics:** Temporal ordering is crucial, sequential patterns
- **Examples:** Stock prices, weather data, sensor readings, website traffic
- **Challenges:** Seasonality, trends, cyclical patterns
- **Algorithms:** ARIMA, LSTM networks, Prophet
- **AWS Services:** Amazon Forecast, Amazon Timestream

Image Data

Definition: Visual data in digital format, represented as pixel values.

- **Formats:** JPEG, PNG, TIFF, raw sensor data

- **Characteristics:** High dimensionality, spatial relationships important
- **Preprocessing:** Resizing, normalization, augmentation
- **Algorithms:** Convolutional Neural Networks (CNNs)
- **AWS Services:** Amazon Rekognition, Amazon SageMaker with computer vision frameworks

Text Data

Definition: Human language in written form, requiring natural language processing techniques.

- **Types:** Documents, social media posts, emails, reviews, transcripts
- **Preprocessing:** Tokenization, stemming, stop-word removal
- **Representation:** Bag of words, TF-IDF, word embeddings, transformers
- **Algorithms:** BERT, GPT, RNNs, transformers
- **AWS Services:** Amazon Comprehend, Amazon Textract, Amazon Bedrock

Structured vs. Unstructured Data

Structured Data

Definition: Data organized in a predefined format with clear schema and relationships.

- **Characteristics:** Fixed fields, consistent format, easily searchable
- **Examples:** Database records, CSV files, XML with schema
- **Storage:** Relational databases, data warehouses
- **Processing:** SQL queries, traditional analytics tools
- **Advantages:** Easy to analyze, query, and process

Unstructured Data

Definition: Data without a predefined structure or organization.

- **Characteristics:** No fixed format, requires preprocessing
- **Examples:** Text documents, images, videos, social media posts, emails
- **Storage:** Data lakes, NoSQL databases, object storage
- **Processing:** Requires specialized tools and techniques
- **Volume:** Represents 80-90% of all data generated
- **AWS Storage:** Amazon S3, Amazon DocumentDB

5. Learning Types in Machine Learning

Supervised Learning

Definition: Learning approach where models are trained on labeled datasets with known input-output pairs.

Process:

1. **Training:** Algorithm learns from labeled examples
2. **Pattern Recognition:** Identifies relationships between inputs and outputs
3. **Prediction:** Applies learned patterns to new, unseen data

4. **Evaluation:** Performance measured against known correct answers

Types:

Classification

- **Goal:** Predict discrete categories or classes
- **Output:** Categorical labels
- **Examples:**
 - Email spam detection (spam/not spam)
 - Image recognition (cat/dog/bird)
 - Medical diagnosis (disease/no disease)
- **Algorithms:** Logistic regression, decision trees, random forest, SVM, neural networks
- **Evaluation Metrics:** Accuracy, precision, recall, F1-score

Regression

- **Goal:** Predict continuous numerical values
- **Output:** Numerical values
- **Examples:**
 - House price prediction
 - Stock price forecasting
 - Temperature prediction
- **Algorithms:** Linear regression, polynomial regression, random forest, neural networks
- **Evaluation Metrics:** Mean Squared Error (MSE), Root Mean Squared Error (RMSE), R-squared

AWS Services: Amazon SageMaker built-in algorithms, Amazon Machine Learning

Unsupervised Learning

Definition: Learning approach where models find hidden patterns in data without labeled examples.

Characteristics:

- **No target variable:** Only input features available
- **Pattern Discovery:** Finds hidden structures in data
- **Exploratory:** Often used for data exploration and understanding

Types:

Clustering

- **Goal:** Group similar data points together
- **Applications:** Customer segmentation, market research, gene sequencing
- **Algorithms:** K-means, hierarchical clustering, DBSCAN
- **Example:** Grouping customers based on purchasing behavior
- **Evaluation:** Silhouette score, within-cluster sum of squares

Association Rules

- **Goal:** Find relationships between different items
- **Applications:** Market basket analysis, recommendation systems
- **Example:** "People who buy bread also buy butter"
- **Algorithms:** Apriori, FP-Growth

Dimensionality Reduction

- **Goal:** Reduce number of features while preserving important information
- **Applications:** Data visualization, noise reduction, feature engineering
- **Algorithms:** Principal Component Analysis (PCA), t-SNE
- **Benefits:** Reduces computational complexity, removes noise

Anomaly Detection

- **Goal:** Identify unusual or outlier data points
- **Applications:** Fraud detection, system monitoring, quality control
- **Algorithms:** Isolation Forest, One-Class SVM, autoencoders

AWS Services: Amazon SageMaker (K-means, PCA), Amazon Fraud Detector

Reinforcement Learning

Definition: Learning approach where an agent learns to make decisions by interacting with an environment and receiving rewards or penalties.

Key Components:

- **Agent:** The learner/decision maker
- **Environment:** The world the agent operates in
- **Actions:** Choices available to the agent
- **States:** Current situation of the agent
- **Rewards:** Feedback signals (positive or negative)
- **Policy:** Strategy for choosing actions

Process:

1. **Agent observes** current state
2. **Selects action** based on current policy
3. **Environment responds** with new state and reward
4. **Agent updates** policy based on reward received
5. **Process repeats** until optimal policy is learned

Learning Strategy:

- **Exploration:** Trying new actions to discover better strategies
- **Exploitation:** Using known good actions to maximize rewards
- **Balance:** Must balance exploration vs exploitation

Applications:

- **Game Playing:** Chess, Go, video games

- **Robotics:** Robot navigation, manipulation tasks
- **Autonomous Systems:** Self-driving cars, drones
- **Resource Management:** Traffic optimization, energy management
- **Finance:** Algorithmic trading, portfolio optimization
- **Personalization:** Content recommendation, ad placement

Algorithms:

- **Q-Learning:** Learns action-value function
- **Policy Gradient:** Directly optimizes policy
- **Actor-Critic:** Combines value and policy methods
- **Deep Q-Networks (DQN):** Q-learning with deep neural networks

AWS Services: AWS DeepRacer, Amazon SageMaker RL

Comparison of Learning Types

Aspect	Supervised	Unsupervised	Reinforcement
Data Type	Labeled examples	Unlabeled data	Interactive environment
Goal	Predict outcomes	Discover patterns	Maximize rewards
Feedback	Correct answers provided	No feedback	Delayed rewards/penalties
Applications	Classification, regression	Clustering, anomaly detection	Game playing, robotics
Evaluation	Accuracy metrics	Internal validation	Reward maximization
Examples	Email spam detection	Customer segmentation	Game AI, autonomous vehicles

Study Tips and Key Takeaways

Memory Aids

1. **AI ⊃ ML ⊃ Deep Learning:** Remember the hierarchical relationship
2. **Supervised = Teacher present, Unsupervised = No teacher, Reinforcement = Learn from consequences**
3. **Batch = Process later, Real-time = Process now**
4. **Structured = Organized, Unstructured = Raw**

Common Exam Focus Areas

- **Definitions:** Know precise definitions of all key terms
- **Relationships:** Understand how AI, ML, and deep learning relate
- **Data types:** Recognize appropriate data types for different problems
- **Learning types:** Match problem scenarios to appropriate learning approaches
- **AWS Services:** Know which AWS services support different AI/ML capabilities

Practice Questions to Consider

1. When would you use batch vs. real-time inferencing?
 2. What type of learning is appropriate for customer segmentation?
 3. How do you address bias in machine learning models?
 4. What's the difference between overfitting and underfitting?
 5. Which AWS services are best for different types of AI workloads?
-

This study guide covers Domain 1, Task Statement 1.1 of the AWS Certified AI Practitioner certification. Focus on understanding concepts rather than memorizing definitions, and practice applying these concepts to real-world scenarios.